

Journal Name

Crossmark

PAPER

RECEIVED
dd Month yyyy
REVISED
dd Month yyyy

Hierarchical Game-Guided Tactical Corridor Planning and Control for Competitive Autonomous Racing

Shuyi Gao, Ren Jin and Defu Lin

School of Aerospace Engineering, Beijing Institute of Technology, Beijing, China

Keywords: Autonomous racing, Tactical corridor planning, Overtaking decision-making.

Abstract

Autonomous racecars in high-speed competitive racing must jointly perform overtaking decision making, trajectory reconstruction, and at-limit control within a confined track space. To address the challenges of highly dynamic vehicle motion, stringent track constraints, complex interactive competition, and low overtaking success rates caused by uncertain feasible overtaking opportunities, this paper proposes a hierarchical game-guided tactical corridor planning and control method. First, a game-based competitive situation model is established in the track Frenet frame under racecar performance-limit constraints to characterize the feasible motion space subject to track boundaries, curvature variations, and tire-friction limits. Then, fixed track priors are embedded into a Truncated Quantile Critics reinforcement learning algorithm at the tactical layer to generate behavior decisions, overtaking-side selections, and time-varying Frenet tactical corridors. Within the tactical corridor, a sparse second-order cone program is employed to generate curvature-regularized paths, while forward-backward speed planning and model predictive contouring control are integrated to achieve dynamically feasible trajectory generation and closed-loop tracking. High-fidelity simulation results show that the proposed method can adaptively reconstruct the overtaking feasible region according to opponent position, track geometry, and vehicle performance boundaries, achieving a 94.7% overtaking success rate over 1000 closed-loop laps.

1 Introduction

With the rapid development of autonomous driving technologies and artificial intelligence methods, autonomous racing has become an important platform for validating high-speed decision making, trajectory planning, and at-limit control of intelligent vehicles. Enhancing the autonomous competitive capability of racecars in adversarial racing scenarios is therefore of great significance. However, high-speed racing is characterized by highly dynamic vehicle motion, strongly constrained track space, and complex interactive competition among vehicles. Under such conditions, merely relying on fixed raceline tracking or single-vehicle trajectory planning makes it difficult to accurately identify feasible overtaking opportunities and achieve stable overtaking. Therefore, developing a hierarchical planning and control method that integrates game-theoretic decision making, trajectory reconstruction, and vehicle dynamic constraints has become a key issue for improving the strategic performance of autonomous racecars in high-speed competitive racing.

To address the coupled challenges of tactical interaction, constrained trajectory generation, and at-limit execution, autonomous racing research has mainly progressed along three directions: behavior planning, local motion planning, and vehicle control. The behavior-planning layer determines high-level racing intentions such as following, defending, and overtaking; the local-planning layer converts these intentions into feasible trajectories within the limited track space; and the control layer ensures accurate execution near the handling limits of the racecar [1, 2, 7]. At the behavior-planning layer, the ego racecar needs to determine its high-level tactical intention, including whether to follow, shadow, defend, or overtake, and on which side the overtaking maneuver should be attempted. Classical autonomous driving frameworks usually rely on rule-based finite-state machines, maneuver libraries, or hierarchical decision modules to select discrete driving behaviors, which are subsequently converted into continuous trajectories by the motion planner [3, 4]. Similar hierarchical ideas have been adopted in autonomous racing systems, where behavior planners switch among raceline following, obstacle avoidance, and overtaking modes according to opponent position, track availability, and racing regulations [5–9]. These rule-based methods are interpretable and easy to integrate into real vehicles, but their tactical

performance strongly depends on manually designed thresholds and track-specific tuning. To explicitly model the strategic coupling between competing vehicles, game-theoretic methods formulate autonomous racing as a non-cooperative dynamic game and compute best-response, Stackelberg, Nash-type, or potential-game strategies [10–16]. Such methods provide a principled way to describe interactive racing behaviors, but exact equilibria of continuous dynamic games are generally difficult to obtain in real time, and approximate solvers may be sensitive to opponent cost assumptions and planning-horizon selection. In parallel, reinforcement learning has shown strong capability in learning complex racing tactics from interaction data [17–21]. However, end-to-end policies usually lack interpretability and cannot directly exploit fixed track priors such as curvature distribution, track width, and vehicle performance boundaries. Therefore, high-speed competitive racing still requires a behavior-planning method that combines learning-based tactical adaptation with structured track and vehicle-limit information.

At the local motion-planning layer, the selected behavior or tactical intention must be transformed into a smooth, collision-free, and dynamically feasible trajectory within the constrained track space. In autonomous racing, many methods are built upon offline raceline optimization, where a minimum-curvature or minimum-time reference line is first computed from the track geometry and vehicle dynamic limits [22–26]. The resulting raceline is effective for single-vehicle lap-time optimization, but competitive scenarios require the ego vehicle to deviate from the nominal raceline and actively reconstruct the overtaking region when an opponent blocks the apex or occupies the preferred racing line. To handle dynamic obstacles and alternative passing opportunities, existing local planners employ Frenet-frame polynomial sampling, lattice planning, graph search, trajectory libraries, homotopy-based trajectory generation, and corridor-constrained optimization [3–5, 27, 28]. These methods can explicitly encode track boundaries, obstacle-avoidance constraints, and part of the vehicle dynamic constraints, which makes them suitable for real-time racing implementation. Nevertheless, many local planners still treat the opponent mainly as a moving obstacle and rely on the behavior planner to provide the passing side or maneuver mode [?, 12]. As a result, the generated trajectory may be locally collision-free but tactically conservative, especially when the available overtaking corridor changes rapidly due to opponent motion and track curvature. Sampling- and graph-based planners may require a large number of candidate trajectories to cover different passing possibilities, while nonlinear optimization methods may suffer from local minima and high computational burden [6, 9]. These limitations motivate a structured local-planning scheme in which the behavior layer first constructs a time-varying tactical corridor, and the local planner then generates a curvature-regularized path and dynamically feasible speed profile inside this corridor [17, 18].

At the control layer, the planned reference must be executed accurately while the vehicle operates close to its handling limits. Model predictive control and model predictive contouring control are widely used in autonomous racing because they jointly handle trajectory tracking, progress maximization, and vehicle dynamic constraints [10, 29–33]; learning-based, robust, and tube-MPC variants further improve accuracy and closed-loop robustness under model mismatch and friction variation [27, 28, 34–36]. However, the control performance depends heavily on the feasibility and smoothness of the planned reference [?, ?, 8]: if the path ignores curvature regularity, speed feasibility, or tire-friction limits, the controller is forced to compensate for dynamically inconsistent commands, leading to oscillatory actuation or degraded robustness [?]. Behavior planning, local motion planning, and vehicle control should therefore be coordinated within a unified hierarchical framework.

Taken together, existing studies have advanced behavior planning, local motion planning, and vehicle control individually, but these modules remain insufficiently coordinated in high-speed competitive racing: behavior planners lack an explicit representation of the feasible overtaking space, local planners have limited tactical modeling of opponent interaction and overtaking timing, and controllers depend on the curvature continuity and tire-friction consistency of the planned reference. It is therefore of great significance to develop a unified hierarchical method that integrates tactical decision making, overtaking feasible-region reconstruction, dynamically feasible trajectory generation, and closed-loop execution, enabling the ego racecar to adaptively reconstruct the overtaking space according to opponent motion, track geometry, and vehicle performance boundaries.

This paper addresses the autonomous racing demand of high-speed competitive autonomous racing, and proposes a hierarchical game-guided tactical corridor planning and control method to solve the problems of overtaking-timing selection, feasible overtaking-space reconstruction, and weak coordination between planning and control under complex interactive racing scenarios. The proposed method embeds fixed track priors into a Truncated Quantile Critics reinforcement

learning policy for tactical decision making, constructs an overtaking feasible-region reconstruction mechanism based on time-varying Frenet tactical corridors, and integrates sparse second-order cone programming, forward–backward speed planning, and model predictive contouring control to achieve dynamically feasible trajectory generation and closed-loop execution. The proposed hierarchical game-guided tactical corridor planning and control method has the following advantages:

1. A fixed-track-prior-guided reinforcement learning method for tactical decision making, which embeds track geometry, curvature, and vehicle performance boundaries into a Truncated Quantile Critics framework to adaptively generate behavior decisions, overtaking-side selections, and time-varying Frenet tactical corridors.
2. A curvature-regularized trajectory planning method under tactical-corridor constraints, where a sparse second-order cone program generates feasible paths and a forward–backward speed strategy enforces vehicle performance limits, connecting feasible-region reconstruction with dynamically feasible trajectory generation.
3. A model predictive contouring control method that jointly regulates contouring error, lag error, inputs, and vehicle dynamic constraints to achieve robust closed-loop tracking of the path–speed reference in complex overtaking scenarios.

The remainder of the paper is organized as follows. Section 2 derives the unified vehicle and Frenet models. Section 3 presents the RL-driven tactical layer. Section 4 details the SOCP path planner and the FB speed profiler. Section 5 formulates the MPCC controller. Section 6 reports two closed-loop overtake simulations. Section 7 concludes.

2 Vehicle, Track and GGV Modeling

2.1 Dynamic Bicycle Model

Let the continuous-time controller state and input be

$$\mathbf{x} = [X, Y, \psi, v_x, v_y, r, \delta, s_p]^\top, \quad \mathbf{u} = [a, \dot{\delta}]^\top, \quad (1)$$

where (X, Y) is the inertial position, ψ is yaw angle, (v_x, v_y) are body-frame longitudinal and lateral velocities, r is yaw rate, δ is front steering angle, and s_p is an auxiliary progress state used only by the MPCC terminal reward. The input a is the commanded longitudinal acceleration and $\dot{\delta}$ is the steering-rate command. The rigid-body kinematics are

$$\dot{X} = v_x \cos \psi - v_y \sin \psi, \quad \dot{Y} = v_x \sin \psi + v_y \cos \psi, \quad \dot{\psi} = r, \quad \dot{s}_p = v_x. \quad (2)$$

The front and rear slip angles are

$$\alpha_f = \arctan\left(\frac{v_y + l_f r}{v_x + \varepsilon_v}\right) - \delta, \quad (3)$$

$$\alpha_r = \arctan\left(\frac{v_y - l_r r}{v_x + \varepsilon_v}\right), \quad (4)$$

where the small positive ε_v regularizes the low-speed denominator. With linear lateral tire stiffness,

$$F_{yf} = -C_f \alpha_f, \quad F_{yr} = -C_r \alpha_r. \quad (5)$$

The resulting Newton–Euler equations are

$$\dot{v}_x = a + r v_y - \frac{1}{m} F_{yf} \sin \delta, \quad (6)$$

$$\dot{v}_y = \frac{1}{m} (F_{yf} \cos \delta + F_{yr}) - r v_x, \quad (7)$$

$$\dot{r} = \frac{1}{I_z} (l_f F_{yf} \cos \delta - l_r F_{yr}). \quad (8)$$

The steering angle evolves as $\dot{\delta} = u_\delta$ where u_δ is the steering-rate input in (1). Equations (2)–(8) define

$$\dot{\mathbf{x}} = f_c(\mathbf{x}, \mathbf{u}), \quad \mathbf{x}_{k+1} = f_d(\mathbf{x}_k, \mathbf{u}_k), \quad (9)$$

where f_d is the fourth-order Runge–Kutta discretization with sampling time T_s . Numerical parameter values are collected in Table 1.




Figure 1: Vehicle and track model in the Frenet frame.

2.2 Track Frenet Coordinates

Let $\mathbf{r}_c(s) \in \mathbb{R}^2$ be the track centerline parameterized by arc length $s \in [0, L_{\text{trk}})$. Its tangent angle is $\theta(s)$ and its curvature is $\kappa(s) = d\theta/ds$. Define the tangent and normal unit vectors as

$$\mathbf{t}(s) = [\cos \theta(s), \sin \theta(s)]^\top, \quad \mathbf{n}_c(s) = [-\sin \theta(s), \cos \theta(s)]^\top. \quad (10)$$

Any vehicle pose sufficiently close to the track is represented by

$$\mathbf{p} = \begin{bmatrix} X \\ Y \end{bmatrix} = \mathbf{r}_c(s) + n \mathbf{n}_c(s), \quad \xi = \psi - \theta(s), \quad (11)$$

where n is the signed lateral offset and ξ is the heading error. The body velocity projected onto the Frenet axes is

$$v_s = v_x \cos \xi - v_y \sin \xi, \quad v_n = v_x \sin \xi + v_y \cos \xi. \quad (12)$$

Differentiating (11) gives the exact kinematic relation

$$\dot{s} = \frac{v_s}{1 - \kappa(s)n}, \quad (13)$$

$$\dot{n} = v_n, \quad (14)$$

$$\dot{\xi} = r - \kappa(s)\dot{s}. \quad (15)$$

The wrapped progress difference used in tactical reasoning is always computed on this same coordinate,

$$\Delta s = ((s_o - s_e + L_{\text{trk}}/2) \bmod L_{\text{trk}}) - L_{\text{trk}}/2, \quad \Delta n = n_o - n_e. \quad (16)$$

For the path planner it is more convenient to use s as the independent variable and to describe a path by $\mathbf{z}(s) = [n(s), \xi(s)]^\top$. Let $\omega(s)$ denote the curvature-like steering variable of this geometric path. The spatial Frenet model is

$$\mathbf{z}'(s) = \mathbf{f}_s(\mathbf{z}(s), \omega(s); s) = \begin{bmatrix} (1 - \kappa n) \tan \xi \\ (1 - \kappa n) \omega / \cos \xi - \kappa \end{bmatrix}. \quad (17)$$

At planning iteration j , this model is linearized about the previous path iterate $(\bar{\mathbf{z}}_k, \bar{\omega}_k)$ at each station s_k ,

$$\mathbf{z}'(s_k) \approx \mathbf{A}_k \mathbf{z}_k + \mathbf{B}_k \omega_k + \mathbf{b}_k, \quad (18)$$

where $\mathbf{A}_k$, $\mathbf{B}_k$ and $\mathbf{b}_k$ are the Jacobian and affine residual of $\mathbf{f}_s$. This single linearized Frenet model is the one used by the SOCP in Sec. 4; no new lateral or heading variables are introduced later.

2.3 GGV Handling Envelope

The speed planner and MPCC share the same handling-limit interpretation. For a speed V , the identified GGV table returns the available tangential and normal acceleration limits $a_T^{\max}(V)$ and $a_N^{\max}(V)$. Combined slip is approximated by the super-elliptic envelope

$$\left| \frac{a_N}{a_N^{\max}(V)} \right|^p + \left| \frac{a_T}{a_T^{\max}(V)} \right|^p \leq 1, \quad (19)$$

with exponent $p = 0.75$. Given a path curvature κ_p and speed V , the normal acceleration is $a_N = \kappa_p V^2$, and the remaining tangential acceleration is

$$a_T^{\text{avail}}(V, a_N) = a_T^{\text{max}}(V) \left[1 - \left(\frac{|a_N|}{a_N^{\text{max}}(V)} \right)^p \right]^{1/p}. \quad (20)$$

Aerodynamic drag is modeled as

$$a_{\text{drag}}(V) = \frac{1}{m} (c_{d,2} V^2 + c_{d,1} V + c_{d,0}), \quad (21)$$

and is added with the acceleration or braking sign convention stated in Sec. 4. The envelope (19) is not a reward term; it is a physical admissibility model used by both the speed profiler and the controller.

3 Tactical Corridor Synthesis

The tactical layer is a learned decision module with a deterministic optimization interface. The reinforcement learning policy outputs a low-dimensional parameter vector encoding the behavior mode, overtaking side, and corridor-shaping parameters, from which a deterministic carver constructs the time-varying Frenet corridor. Let $\mathcal{B}(s) = [n_{\text{trk}}^-(s), n_{\text{trk}}^+(s)]$ denote the track boundary in the Frenet coordinates of Sec. 2, and let $\mathcal{O}_t$ denote the observed opponent set. Online operation is written as the composition

$$\begin{aligned} \tilde{\mathbf{o}}_t &= \mathcal{E}(\mathbf{x}_{e,t}, \mathcal{O}_t, \mathcal{P}, m_t, \sigma_t^{\text{lock}}), \\ \mathbf{u}_t &\sim \pi_\theta(\cdot | \tilde{\mathbf{o}}_t), \\ \mathbf{a}_t &= \Pi_{\mathcal{A}} \left(\underline{\mathbf{a}} + \frac{\mathbf{u}_t + \mathbf{1}}{2} \odot (\bar{\mathbf{a}} - \underline{\mathbf{a}}) \right), \\ m_{t+1} &= \mathcal{F}_m(\mathbf{o}_t, \mathbf{a}_t, m_t, \sigma_t^{\text{lock}}), \\ \mathcal{C}_t &= \mathcal{F}_c(\mathcal{B}, \mathcal{O}_t, \mathbf{a}_t, m_{t+1}, \sigma_t^{\text{lock}}), \\ \mathcal{G}_t &= (m_{t+1}, V_{\text{cap},t}, \mathcal{C}_t, \pi_t). \end{aligned} \quad (22)$$

Only π_θ is learned. The projection $\Pi_{\mathcal{A}}$, the finite-state machine $\mathcal{F}_m$, and the corridor carver $\mathcal{F}_c$ are deterministic maps that convert the policy output into auditable constraints for the planner. The training problem is

$$\pi_\theta^* = \arg \max_{\pi_\theta} \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^{T_{\text{ep}}-1} \gamma^t r(\tilde{\mathbf{o}}_t, \mathbf{a}_t, m_t, \mathcal{C}_t) \right], \quad (23)$$

so reinforcement learning is used to select a tactical feasible set rather than low-level steering or throttle.

The encoder $\mathcal{E}$ combines a duel-state observation, a fixed track prior, and a vehicle-capability prior. The dynamic observation is

$$\mathbf{o}_t^{\text{ego}} = [\bar{V}_e, \bar{\xi}_e, \bar{n}_e, \bar{a}_{T,e}, \bar{a}_{N,e}, \bar{s}_e]^\top, \quad (24)$$

$$\mathbf{o}_t^{\text{opp}} = [\bar{\Delta}s, \bar{\Delta}n, \bar{\Delta}V_s, \bar{s}_o, \bar{n}_o, \bar{g}, \bar{\tau}_{\text{tc}}, \bar{V}_o]^\top, \quad (25)$$

$$\mathbf{o}_t = [\mathbf{o}_t^{\text{ego}}, \mathbf{o}_t^{\text{opp}}, \mathbf{b}_t, \sigma_t^{\text{lock}}, \mathbf{e}_{m_t}]. \quad (26)$$

With five FSM modes, $\mathbf{o}_t \in \mathbb{R}^{23}$. Each scalar is normalized as

$$\bar{z} = \text{sat}(z/z_{\text{char}}), \quad \text{sat}(q) = \min\{1, \max\{-1, q\}\}, \quad (27)$$

using the characteristic values in Table 1. The relative progress and lateral separation are given by (16), $\Delta V_s = V_o - V_e$, and $g = \sqrt{\Delta s^2 + \Delta n^2}$. The clipped time-to-collision feature is

$$\tau_{\text{tc}} = \begin{cases} \text{clip}(\Delta s / \max(-\Delta V_s, \varepsilon_\tau), 0, \tau_{\text{tc},\text{max}}), & \Delta s > 0, \Delta V_s < 0, \\ \tau_{\text{tc},\text{max}}, & \text{otherwise,} \end{cases} \quad (28)$$

so only a closing leader creates a finite urgency scale. The bucket vector $\mathbf{b}_t$ coarsely encodes whether the ego is fast-closing, speed-matched or receding. The memory variables σ_t^{lock} and $\mathbf{e}_{m_t}$ make side commitment and hysteresis observable to the policy.

The dynamic duel state is augmented by a non-trainable map $\mathcal{P} = \{\rho_i\}_{i=0}^{M-1}$. This map contains not only curvature and width, but also a local estimate of vehicle capability. From the GGV envelope in Sec. 2, define the curvature-limited speed preview

$$V_{\text{ggv}}(s) = \sup_{0 \leq V \leq v_{x,\max}} \{V : |\kappa(s)|V^2 \leq a_N^{\max}(V)\}, \quad (29)$$

with $V_{\text{ggv}} = v_{x,\max}$ when the curvature constraint is inactive. For a preview interval $\mathcal{I}(s; d_a, d_b) = [s + d_a, s + d_b]$, the statistics inserted into the policy are

$$\kappa_{d_a:d_b}^{\max}(s) = \max_{\ell \in \mathcal{I}(s; d_a, d_b)} |\kappa(\ell)|, \quad W_{\min}^{L/R}(s) = \min_{\ell \in [s, s+d_p]} |n_{\text{trk}}^{+/-}(\ell)|, \quad V_{\min}^{\text{ggv}}(s) = \min_{\ell \in [s, s+d_p]} V_{\text{ggv}}(\ell). \quad (30)$$

Querying $\mathcal{P}$ at progress s returns a 14-dimensional vector

$$\rho(s) = \underbrace{[\bar{\kappa}_0, \bar{\kappa}_{0:d_1}^{\max}, \dots, \bar{\kappa}_{d_3:d_p}^{\max}, \bar{\kappa}_{\text{mean}}, \bar{\kappa}_{\text{std}}]}_{\text{curvature preview (7)}}, \underbrace{[\bar{W}_{\min}^L, \bar{W}_{\min}^R, \bar{W}_0^L, \bar{W}_0^R, \Delta \bar{W}_{\min}]}_{\text{width (5)}}, \underbrace{[\bar{V}_0^{\text{ggv}}, \bar{V}_{\min}^{\text{ggv}}]}_{\text{GGV (2)}}^{\top}. \quad (31)$$

The curvature terms encode upcoming cornering severity at four look-ahead intervals, the width terms encode whether an overtake side remains geometrically open, and the GGV terms encode whether the vehicle can exploit that side without exceeding lateral-acceleration capability. The final actor-critic input is

$$\tilde{\mathbf{o}}_t = [\mathbf{o}_t, \rho(s_e), \rho(s_o)] \in \mathbb{R}^{51}, \quad (32)$$

so an identical relative gap can yield different actions on a straight, at corner entry, or before a narrow exit.

The actor output is not a trajectory but a ten-dimensional tactical parameter vector. A sample $\mathbf{u}_t \in [-1, 1]^{10}$ is mapped to

$$\mathbf{a}_t = \Pi_{\mathcal{A}} \left(\underline{\mathbf{a}} + \frac{1}{2}(\mathbf{u}_t + \mathbf{1}) \odot (\bar{\mathbf{a}} - \underline{\mathbf{a}}) \right), \quad (33)$$

$$\mathbf{a}_t = [L_t^{\text{cap}}, R_t^{\text{cap}}, d_{\text{safe}}, V_{\text{cap}}, g_{\text{chase}}, g_{\text{ot}}, g_{\text{abort}}, b_{\text{lat}}, w_{\text{safe}}, q_{\text{side}}]^{\top}. \quad (34)$$

The projection $\Pi_{\mathcal{A}}$ enforces

$$\mathcal{A} = \{\mathbf{a} : g_{\text{chase}} \geq g_{\text{ot}} + \varepsilon_{co}, g_{\text{abort}} \geq g_{\text{chase}} + \varepsilon_{ca}, L_t^{\text{cap}} + R_t^{\text{cap}} \geq \varepsilon_{\text{corr}}\}. \quad (35)$$

The components of $\mathbf{a}_t$ have direct downstream meanings: $(L_t^{\text{cap}}, R_t^{\text{cap}})$ cap the left/right corridor width, d_{safe} sets opponent clearance, V_{cap} caps the speed profile, $(g_{\text{chase}}, g_{\text{ot}}, g_{\text{abort}})$ tune FSM hysteresis, b_{lat} sets the far-field lateral bias, w_{safe} controls recovery width, and q_{side} chooses the pass side,

$$\sigma_t = \begin{cases} +1, & q_{\text{side}} \geq 0 \quad \text{left-side pass,} \\ -1, & q_{\text{side}} < 0 \quad \text{right-side pass.} \end{cases} \quad (36)$$

Once σ_t is accepted in OVERTAKE, it is copied into σ_t^{lock} and held until the duel terminates.

The continuous tactical action is converted into a persistent intention by a hysteretic finite-state machine with modes $\mathcal{M} = \{\text{RACELINE}, \text{SHADOW}, \text{OVERTAKE}, \text{HOLD}, \text{DEFEND}\}$. Let g_{rear} denote the closest rear gap and $\bar{c}_t$ indicate that the post-overtake immunity window has expired. The selected side is considered ready when the previewed side width and GGV speed are both admissible,

$$\text{ready}_t(\sigma) = \mathbb{1} \left[\min_{s \in [s_e, s_e+d_p]} W_{\sigma}(s) \geq L_{\min}^{\text{corr}} \right] \mathbb{1} \left[\min_{s \in [s_e, s_e+d_p]} V_{\text{ggv}}(s) \geq \min\{V_e, V_{\text{cap}}\} \right], \quad (37)$$

where W_{+1} and W_{-1} are the left/right widths remaining after opponent exclusion. The mode update is

$$m_{t+1} = \begin{cases} \text{OVERTAKE,} & g \leq g_{\text{ot}} \wedge \text{ready}_t(\sigma_t) \wedge \bar{c}_t, \\ \text{SHADOW,} & g_{\text{ot}} < g \leq g_{\text{chase}}, \\ \text{HOLD,} & m_t = \text{OVERTAKE} \wedge \text{abort}_t, \\ \text{DEFEND,} & g > g_{\text{def}}^{\text{front}} \wedge 0 < g_{\text{rear}} < g_{\text{def}}^{\text{rear}} \wedge \bar{c}_t, \\ \text{RACELINE,} & g > g_{\text{chase}}, \\ m_t, & \text{otherwise.} \end{cases} \quad (38)$$




Figure 2: Geometric semantics of mode-dependent tactical corridors. The feasible corridor is released to the full track width in RACELINE/HOLD, funneled toward the chosen passing side behind the opponent in SHADOW, reshaped for side-specific passing in OVERTAKE, and placed in front of the following opponent in DEFEND.

The abort flag is

$$\text{abort}_t = (g > g_{\text{abort}}) \vee (\Delta n_{\min} < \Delta n_{\text{safe}} \wedge |\Delta s| < g_{\text{interlock}}), \quad (39)$$

where $\Delta n_{\min}$ is the minimum lateral separation predicted within the local interaction window. Therefore the policy adjusts thresholds and side preference, while the hybrid logic preserves a persistent mode sequence for the downstream optimizer.

Figure 2 illustrates the geometric meaning of the FSM modes and the corresponding tactical corridors. The green shaded area denotes the feasible corridor published by the tactical layer, and the green curves denote its lateral boundaries. RACELINE and HOLD share the same full-track-width corridor; the former encourages lap-time-oriented driving, whereas the latter keeps the corridor open but regulates the speed profile through $V_{\text{cap},t}$ to maintain a safe gap to the preceding vehicle. SHADOW, OVERTAKE, and DEFEND further reshape the corridor according to the opponent position, respectively guiding the ego to the rear-side of the opponent on the chosen passing side in preparation for the pass, supporting side-specific passing, and blocking the opponent's forward driving direction.

The reward is defined on tactical outcome and constraint quality. At each step,

$$r_t = w_{\text{ot}}\phi_{\text{ot}} + w_{\text{ovb}}\phi_{\text{ovb}} + w_{\text{col}}\phi_{\text{col}} + w_{\text{lat}}\phi_{\text{lat}} + w_{\text{crush}}\phi_{\text{crush}} + w_{\text{app}}\phi_{\text{app}} + w_{\text{sm}}\phi_{\text{sm}} + w_{\text{gs}}\phi_{\text{gs}} + w_{\text{sbs}}\phi_{\text{sbs}} + w_{\text{wide}}\phi_{\text{wide}}, \quad (40)$$

with weights listed in Table 1. The event and geometry features are

$$\phi_{\text{ot}} = \mathbb{K}[\Delta s_{t-1} > 0] \mathbb{K}[\Delta s_t < 0] \mathbb{K}[\bar{c}_t], \quad (41)$$

$$\phi_{\text{ovb}} = \mathbb{K}[\Delta s_{t-1} < 0] \mathbb{K}[\Delta s_t > 0], \quad (42)$$

$$\phi_{\text{col}} = \mathbb{K}[\text{contact} \vee \text{offtrack} \vee \text{comm.loss}], \quad (43)$$

$$\phi_{\text{lat}} = \mathbb{K}[m_t = \text{RACELINE}] \left(\frac{n_e - n_{\text{rl}}(s_e)}{n_{\text{char}}} \right)^2, \quad (44)$$

$$\phi_{\text{crush}} = \mathbb{K}[|\Delta s| < g_{\text{interlock}}] \left(\frac{\max\{0, \Delta n_{\text{safe}} - |\Delta n|\}}{n_{\text{char}}} \right)^2. \quad (45)$$

The approach feature is

$$\phi_{\text{app}} = \mathbb{K}[\Delta s > 0] \mathbb{K}[g_t < g_{\text{chase}}] \text{clip}\left(\frac{g_{t-1} - g_t}{g_{\text{char}}}, 0, 1\right), \quad (46)$$

and the regularizers applied to the policy output and the published corridor are

$$\phi_{\text{sm}} = \|\mathbf{a}_t - \mathbf{a}_{t-1}\|_{\mathbf{W}_a}^2, \quad (47)$$

$$\phi_{\text{gs}} = \frac{1}{N_{\text{sta}}} \sum_{k=0}^{N_{\text{sta}}-1} \left[(n_{k,t}^+ - n_{k,t-1}^+)^2 + (n_{k,t}^- - n_{k,t-1}^-)^2 \right], \quad (48)$$

$$\phi_{\text{sbs}} = \mathbb{I}[\Delta s < g_{\text{interlock}}] \mathbb{I}[|\Delta n| \geq \Delta n_{\text{safe}}], \quad (49)$$

$$\phi_{\text{wide}} = \frac{1}{N_{\text{sta}}} \sum_{k=0}^{N_{\text{sta}}-1} \text{clip}\left(\frac{n_k^+ - n_k^- - L_{\text{min}}^{\text{corr}}}{n_{\text{char}}}, 0, 1\right). \quad (50)$$

Positive weights reward pass completion, useful gap closing, safe side-by-side tenure and non-degenerate corridors. Negative weights penalize being overtaken, collision, unnecessary raceline deviation, lateral compression, action jitter and corridor jitter. Hence the reward shapes only the tactical interface; it does not modify the SOCP or MPCC objectives.

The policy is optimized with Truncated Quantile Critics (TQC), a distributional off-policy algorithm chosen specifically because its pessimistic value estimate reduces the risk of committing to corridors that appear advantageous in expectation but carry high collision probability in the tail—a property validated by the ablation study in Sec. 6 where replacing TQC with standard SAC increases collisions from zero to five. Two critics parameterized by φ_i approximate the return distribution with N_q quantile atoms $\{z_{\varphi_i}^j(\tilde{\mathbf{o}}, \mathbf{u})\}_{j=1}^{N_q}$. Given a replay transition, draw $\mathbf{u}' \sim \pi_{\theta}(\cdot|\tilde{\mathbf{o}}')$ and form the truncated target

$$\mathcal{Z}_{\text{targ}} = \text{sort}\left(\{r_t + \gamma(1 - d_t)[z_{\varphi_i}^j(\tilde{\mathbf{o}}', \mathbf{u}') - \alpha \log \pi_{\theta}(\mathbf{u}'|\tilde{\mathbf{o}}')]\}_{i,j}\right)_{1:K}. \quad (51)$$

Retaining only the lowest K out of $2N_q$ atoms suppresses overestimation of rare high-reward overtakes that coincide with near-miss geometries. The critic minimizes the quantile Huber loss

$$J_Q(\varphi) = \sum_{i=1}^2 \sum_{j=1}^{N_q} \mathbb{E}_{\mathcal{D}} \left[\rho_{\tilde{\tau}_j}^{\kappa_q} (y - z_{\varphi_i}^j(\tilde{\mathbf{o}}, \mathbf{u}_t)) \right], \quad (52)$$

where $y \in \mathcal{Z}_{\text{targ}}$, and the actor maximizes the truncated Q-value with entropy regularization:

$$J_{\pi}(\theta) = \mathbb{E}_{\tilde{\mathbf{o}} \sim \mathcal{D}, \mathbf{u} \sim \pi_{\theta}} \left[\alpha \log \pi_{\theta}(\mathbf{u}|\tilde{\mathbf{o}}) - \frac{1}{K} \sum_{z \in \mathcal{Z}_K(\tilde{\mathbf{o}}, \mathbf{u})} z \right]. \quad (53)$$

The entropy temperature α is tuned automatically. Training episodes randomize initial gap, relative speed, lateral offset and opponent behavior; termination occurs at pass completion, being overtaken, collision, off-track excursion, communication loss or timeout. Evaluation uses the deterministic mean action with the same FSM and corridor carver.

The tactical output consumed by the planner is

$$\mathcal{G}_t = (m_t, V_{\text{cap},t}, \mathcal{C}_t, \pi_t), \quad \mathcal{C}_t = \{(s_k, n_k^-, n_k^+)\}_{k=0}^{N_{\text{sta}}-1}, \quad (54)$$

where π_t stores race-control permission. The carver is the deterministic map

$$\mathcal{C}_t = \mathcal{F}_c \left(\{n_{\text{trk}}^-(s_k), n_{\text{trk}}^+(s_k)\}_{k=0}^{N_{\text{sta}}-1}, \mathcal{O}_t, m_t, \sigma_t^{\text{lock}}, \mathbf{a}_t \right). \quad (55)$$

It first initializes the station grid and physical bounds,

$$s_k = s_e + k\Delta s_{\text{sta}}, \quad n_k^- = n_{\text{trk}}^-(s_k), \quad n_k^+ = n_{\text{trk}}^+(s_k). \quad (56)$$

For an opponent with heading error ξ_o and footprint (L_v, W_v) , the half-extents projected onto Frenet axes are

$$\rho_n = \frac{1}{2}(L_v |\sin \xi_o| + W_v |\cos \xi_o|), \quad \rho_s = \frac{1}{2}(L_v |\cos \xi_o| + W_v |\sin \xi_o|), \quad (57)$$

and the lateral exclusion radius is

$$\rho_{\text{exc}} = \rho_n + \frac{W_v}{2} + d_{\text{safe}}. \quad (58)$$

The longitudinal influence of the footprint is

$$\phi_{\text{fade}}(\Delta s_k) = \begin{cases} \cos^{p_{\text{fade}}} \left(\frac{\pi}{2} \frac{|\Delta s_k|}{S_{\text{fade}}} \right), & |\Delta s_k| < S_{\text{fade}}, \\ 0, & \text{otherwise,} \end{cases} \quad (59)$$

multiplied by a startup ramp $\phi_{\text{sr}}(k)$ and a distance ramp $\phi_{\text{rmp}}(k)$. In OVERTAKE, the bound on the opponent side is tightened and the selected side is left open:

$$\begin{cases} n_k^- \leftarrow \max(n_k^-, n_{o,k} + \rho_{\text{exc}} \phi_{\text{fade}} \phi_{\text{sr}} \phi_{\text{rmp}}), & \sigma_t = +1, \\ n_k^+ \leftarrow \min(n_k^+, n_{o,k} - \rho_{\text{exc}} \phi_{\text{fade}} \phi_{\text{sr}} \phi_{\text{rmp}}), & \sigma_t = -1. \end{cases} \quad (60)$$

In SHADOW, the interval is shaped into an asymmetric funnel toward σ_t ; in HOLD, it is relaxed toward the raceline; in DEFEND, it is biased toward the defensive side. The learned bias and width caps are then blended into the far field. With

$$w_k = \begin{cases} (k/k_{\text{blend}})^2, & k < k_{\text{blend}}, \\ 1, & k \geq k_{\text{blend}}, \end{cases} \quad (61)$$

$$n_k^+ \leftarrow (1 - w_k)n_k^+ + w_k \min(n_k^+, b_{\text{lat}} + L_t^{\text{cap}}), \quad (62)$$

$$n_k^- \leftarrow (1 - w_k)n_k^- + w_k \max(n_k^-, b_{\text{lat}} - R_t^{\text{cap}}). \quad (63)$$

The minimum-width guard is

$$n_k^+ - n_k^- \geq L_{\text{min}}^{\text{corr}}. \quad (64)$$

If this inequality cannot be recovered within the physical track boundary, the previous feasible interval is held at that station. Finally, the published corridor is temporally filtered,

$$\begin{bmatrix} n_k^- \\ n_k^+ \end{bmatrix}_t \leftarrow \beta \begin{bmatrix} n_k^- \\ n_k^+ \end{bmatrix}_t + (1 - \beta) \begin{bmatrix} n_k^- \\ n_k^+ \end{bmatrix}_{t-1}, \quad (65)$$

and the tuple $\mathcal{G}_t$ is sent to the SOCP planner. Thus the learned layer contributes exactly three planning quantities: m_t selects the mode-dependent terminal target, $V_{\text{cap},t}$ clips the GGV speed profile, and $\mathcal{C}_t$ supplies the lateral inequalities $n_k^- \leq n_k \leq n_k^+$ in Sec. 4. If a tactical update is delayed by the simulation interface, the last valid $\mathcal{G}_t$ is held for one planning cycle.

4 Local Planning

The local planner converts the tactical tuple $\mathcal{G}_t = (m_t, V_{\text{cap},t}, \mathcal{C}_t, \pi_t)$ into a time-parameterized reference. Its structure is deliberately split:

$$\mathcal{G}_t \xrightarrow{\text{ECOS SOCP}} \{s_k, n_k^*, \xi_k^*, \kappa_{p,k}^*\}_{k=0}^{N-1} \xrightarrow{\text{FB speed surface}} \mathcal{R}_t. \quad (66)$$

The first stage solves a sparse second-order cone program for the Frenet path. The second stage computes a forward-backward (FB) speed surface on the resulting local path using the GGV envelope of Sec. 2.

The horizon is sampled at

$$s_k = s_e + k\Delta s_{\text{sta}}, \quad k = 0, \dots, N-1, \quad (67)$$

where the wrapped value of s_k is used when querying track data. The tactical corridor is intersected with the physical track boundary,

$$\underline{n}_k = \max\{n_{\text{trk}}^-(s_k), n_k^-\}, \quad \bar{n}_k = \min\{n_{\text{trk}}^+(s_k), n_k^+\}. \quad (68)$$

If no tactical corridor is active, $\underline{n}_k$ and $\bar{n}_k$ are the track limits after the standard safety margin. The optimization variables are

$$\mathbf{y} = [\{n_k, \xi_k, \omega_k, \sigma_k^\omega\}_{k=0}^{N-1}, \epsilon_n, \epsilon_\xi], \quad \mathbf{z}_k = \begin{bmatrix} n_k \\ \xi_k \end{bmatrix}, \quad (69)$$

where (n_k, ξ_k) are the Frenet variables of Sec. 2, ω_k is the curvature-like path input in (17), σ_k^ω is the cone epigraph for ω_k^2 , and $(\epsilon_n, \epsilon_\xi)$ are terminal relaxation variables.

The initial condition is hard,

$$\mathbf{z}_0 = \begin{bmatrix} n_e \\ \xi_e \end{bmatrix}, \quad (70)$$

while the terminal target is mode dependent,

$$n_T = \begin{cases} n_{\text{rl}}(s_{N-1}), & m_t = \text{RACELINE}, \quad n_{\text{rl}}(s_{N-1}) \in [\underline{n}_{N-1}, \bar{n}_{N-1}], \\ \frac{1}{2}(\underline{n}_{N-1} + \bar{n}_{N-1}), & \text{otherwise}, \end{cases} \quad \xi_T = 0. \quad (71)$$

Instead of forcing this terminal condition as an equality, the implemented planner uses soft terminal inequalities,

$$-\epsilon_n \leq n_{N-1} - n_T \leq \epsilon_n, \quad -\epsilon_\xi \leq \xi_{N-1} - \xi_T \leq \epsilon_\xi, \quad \epsilon_n \geq 0, \quad \epsilon_\xi \geq 0. \quad (72)$$

This prevents ECOS infeasibility when a short horizon begins inside a tight overtake corridor but still biases the path toward the tactical terminal intent.

At sequential-linearization iteration j , the nonlinear spatial model (17) is linearized at $(\mathbf{z}_k^j, \omega_k^j)$,

$$\mathbf{f}_s(\mathbf{z}_k, \omega_k; s_k) \approx \mathbf{A}_k^j \mathbf{z}_k + \mathbf{B}_k^j \omega_k + \mathbf{b}_k^j \triangleq \hat{\mathbf{f}}_{s,k}^j(\mathbf{z}_k, \omega_k). \quad (73)$$

Trapezoidal integration gives the linear equality constraint

$$\mathbf{z}_{k+1} - \mathbf{z}_k = \frac{\Delta s_k}{2} \left[\hat{\mathbf{f}}_{s,k}^j(\mathbf{z}_k, \omega_k) + \hat{\mathbf{f}}_{s,k+1}^j(\mathbf{z}_{k+1}, \omega_{k+1}) \right], \quad k = 0, \dots, N-2. \quad (74)$$

The lateral corridor and curvature epigraph are

$$\underline{n}_k \leq n_k \leq \bar{n}_k, \quad k = 0, \dots, N-1, \quad (75)$$

$$\left\| \begin{bmatrix} 2\omega_k \\ \sigma_k^\omega - 1 \end{bmatrix} \right\|_2 \leq \sigma_k^\omega + 1, \quad k = 0, \dots, N-1, \quad (76)$$

which is equivalent to $(\omega_k)^2 \leq \sigma_k^\omega$ with $\sigma_k^\omega \geq 0$. The SOCP solved by ECOS is

$$\begin{aligned} \mathbf{y}^{j+1} = \arg \min_{\mathbf{y}} \quad & \sum_{k=0}^{N-1} w_\omega \sigma_k^\omega + w_{n,T} \epsilon_n + w_{\xi,T} \epsilon_\xi \\ \text{s.t.} \quad & (70), (72), (74), (75), \text{ and } (76). \end{aligned} \quad (77)$$

Equivalently, ECOS receives the conic form

$$\min_{\mathbf{y}} \mathbf{c}^\top \mathbf{y}, \quad \mathbf{A}_{\text{eq}} \mathbf{y} = \mathbf{b}_{\text{eq}}, \quad \mathbf{G}_{\text{lp}} \mathbf{y} \leq \mathbf{h}_{\text{lp}}, \quad \|\mathbf{G}_{q,k} \mathbf{y} + \mathbf{g}_{q,k}\|_2 \leq \mathbf{h}_{q,k}^\top \mathbf{y} + h_{q,k}. \quad (78)$$

After each solve, the linearization trajectory is updated. The loop stops when

$$\Delta_n^j = \max_k |n_k^{j+1} - n_k^j| \leq \varepsilon_n^{\text{sl}}, \quad \Delta_\xi^j = \max_k |\xi_k^{j+1} - \xi_k^j| \leq \varepsilon_\xi^{\text{sl}}, \quad (79)$$

or when the iteration limit is reached. The optimized Cartesian path and its heading are

$$\mathbf{p}_k^* = \mathbf{r}_c(s_k) + n_k^* \mathbf{n}_c(s_k), \quad \psi_k^* = \theta(s_k) + \xi_k^*, \quad \kappa_{p,k}^* = \omega_k^*. \quad (80)$$

The path arclength used by the speed planner is then recomputed from the optimized Cartesian points,

$$\ell_0 = 0, \quad \ell_{k+1} = \ell_k + \|\mathbf{p}_{k+1}^* - \mathbf{p}_k^*\|_2, \quad \Delta \ell_k = \ell_{k+1} - \ell_k. \quad (81)$$

The FB speed profiler computes the largest speed surface that remains consistent with curvature, longitudinal acceleration and braking. First, a curvature speed ceiling is obtained at each station:

$$V_{\text{lim},k} = \begin{cases} v_{x,\text{max}}, & |\kappa_{p,k}^*| < \varepsilon_\kappa, \\ \lim_{r \rightarrow \infty} V_k^{(r)}, & V_k^{(r+1)} = \min \left\{ v_{x,\text{max}}, \sqrt{\frac{a_N^{\text{max}}(V_k^{(r)})}{|\kappa_{p,k}^*|}} \right\}, \end{cases} \quad (82)$$

initialized with $V_k^{(0)} = v_{x,\text{max}}$. For any candidate speed V at station k , define

$$a_{N,k}(V) = |\kappa_{p,k}^*| V^2, \quad \chi_k(V) = \left[1 - \left(\frac{a_{N,k}(V)}{a_N^{\text{max}}(V)} \right)^p \right]_+^{1/p}, \quad (83)$$

where $[\cdot]_+ = \max\{0, \cdot\}$ and p is the GGV exponent in (19). The available forward and braking accelerations are

$$a_{T,k}^+(V) = a_{T,k}^{+,max}(V)\chi_k(V) + a_{drag}(V), \quad (84)$$

$$a_{T,k}^-(V) = a_{T,k}^{-,max}(V)\chi_k(V) - a_{drag}(V), \quad (85)$$

where $a_{T,k}^{+,max}$ and $a_{T,k}^{-,max}$ are the calibrated driving and braking envelopes. Under the sign convention of (21), a_{drag} is usually negative at racing speeds; hence it reduces forward acceleration and assists braking.

The forward sweep starts from the measured speed,

$$V_0^f = \min\{\max(V_e, V_{min}), V_{lim,0}\}, \quad (86)$$

and propagates

$$V_{k+1}^f = \text{clip}_{[V_{min}, V_{lim,k+1}]} \left(\sqrt{(V_k^f)^2 + 2\Delta\ell_k a_{T,k}^+(V_k^f)} \right), \quad k = 0, \dots, N-2. \quad (87)$$

The backward sweep is initialized by

$$V_{N-1}^b = \begin{cases} \min\{V_T, V_{lim,N-1}\}, & \text{if a terminal speed is enforced,} \\ \min\{V_{N-1}^f, V_{lim,N-1}\}, & \text{otherwise,} \end{cases} \quad (88)$$

and propagates

$$V_k^b = \text{clip}_{[V_{min}, V_{lim,k}]} \left(\sqrt{(V_{k+1}^b)^2 + 2\Delta\ell_k a_{T,k+1}^-(V_{k+1}^b)} \right), \quad k = N-2, \dots, 0. \quad (89)$$

The final speed surface is the pointwise feasible envelope

$$V_k^* = \min\{V_k^f, V_k^b, V_{lim,k}, V_{cap,t}\}. \quad (90)$$

The accelerations stored with the local reference are

$$a_{T,k}^* = \frac{(V_{k+1}^*)^2 - (V_k^*)^2}{2\Delta\ell_k}, \quad a_{N,k}^* = \kappa_{p,k}^* (V_k^*)^2. \quad (91)$$

Finally, the time stamps are reconstructed by trapezoidal time integration,

$$\tau_0^r = 0, \quad \tau_{k+1}^r = \tau_k^r + \frac{2\Delta\ell_k}{V_k^* + V_{k+1}^* + \varepsilon_v}, \quad (92)$$

and the planner publishes

$$\mathcal{R}_t = \{(X_k^r, Y_k^r, \psi_k^r, \kappa_{p,k}^r, V_k^r, \tau_k^r)\}_{k=0}^{N_{out}-1}, \quad (93)$$

after resampling the path, speed and time arrays to the controller cadence.

5 Model Predictive Contouring Control

The control layer tracks the path-speed reference $\mathcal{R}_t$ from (93) by solving a receding-horizon nonlinear program at the actuation rate $1/T_c$ ($T_c = T_s = 0.05$ s). A standard path-tracking MPC penalizes Cartesian deviations $(x - x^r)^2 + (y - y^r)^2$; this couples longitudinal timing to lateral accuracy and causes aggressive re-plans when the tactical layer shifts the corridor. By contrast, the MPCC formulation decomposes the error into a *contouring* component (normal to the path) and a *lag* component (tangential to the path), and penalizes only the former. Longitudinal timing is recovered through a speed-reference term, so a corridor-side switch only changes the reference sequence without requiring controller re-initialization.

At control instant t the current vehicle pose defines the local origin. The state is the planar pose $\mathbf{x}_k = [x_k, y_k, \theta_k]^\top$ and the input is the commanded speed and front-wheel steering angle $\mathbf{v}_k = [V_k, \delta_k]^\top$. The initial condition is

$$\mathbf{x}_{0|t} = \mathbf{0}, \quad \mathbf{v}_{-1|t} = [V_{-1|t}, \delta_{-1|t}]^\top, \quad (94)$$

where $\boldsymbol{\nu}_{-1|t}$ is the command applied in the previous cycle, retained for rate-constraint enforcement. The prediction model is a kinematic bicycle with an identified nonlinear yaw-rate map,

$$\boldsymbol{\chi}_{k+1|t} = \boldsymbol{\chi}_{k|t} + T_c \begin{bmatrix} V_{k|t} \cos \theta_{k|t} \\ V_{k|t} \sin \theta_{k|t} \\ \Omega(V_{k|t}, \delta_{k|t}) \end{bmatrix}, \quad (95)$$

$$\Omega(V, \delta) = \frac{\tan(a_\delta \delta + b_\delta)}{L_c} (c_V V + d_V), \quad (96)$$

where L_c is the wheelbase and $\{a_\delta, b_\delta, c_V, d_V\}$ are identified from closed-loop vehicle data using least-squares regression on measured yaw rate, speed, and steering angle. The model is deliberately low-dimensional: the GGV speed profiler in Sec. 4 has already ensured that the reference speed is dynamically feasible, so the controller need only reject pose disturbances at the actuation rate $1/T_c$.

The reference sequence is obtained by transforming the planner output into the local frame:

$$\bar{\mathbf{r}}_{k|t} = [x_k^r, y_k^r, \theta_k^r, V_k^r, \delta_k^r]^\top, \quad k = 1, \dots, N_c, \quad (97)$$

where the reference steering is reconstructed from the planned path curvature $\kappa_{p,k}^r$ via the inverse kinematic relation

$$\delta_k^r = \arctan\left(\frac{L_c \kappa_{p,k}^r V_k^r}{V_k^r + \varepsilon_v}\right). \quad (98)$$

For each prediction step, define the position error in the moving tangent-normal frame of the reference path. Let $\mathbf{t}_k^r = [\cos \theta_k^r, \sin \theta_k^r]^\top$ and $\mathbf{n}_k^r = [-\sin \theta_k^r, \cos \theta_k^r]^\top$ be the unit tangent and normal at step k . The contouring error (normal deviation), lag error (tangential deviation), heading error, and input-tracking error are

$$e_{c,k|t} = (\mathbf{n}_{k|t}^r)^\top ([x_{k|t} - x_k^r, y_{k|t} - y_k^r]^\top), \quad (99)$$

$$e_{\ell,k|t} = (\mathbf{t}_{k|t}^r)^\top ([x_{k|t} - x_k^r, y_{k|t} - y_k^r]^\top), \quad (100)$$

$$e_{\psi,k|t} = \text{wrap}(\theta_{k|t} - \theta_k^r), \quad (101)$$

$$\mathbf{e}_{\nu,k|t} = \boldsymbol{\nu}_{k|t} - \boldsymbol{\nu}_k^r. \quad (102)$$

The contouring error e_c measures the lateral deviation from the planned path and is the primary tracking objective. The lag error e_ℓ is not penalized in the cost; longitudinal timing is handled implicitly through the speed reference in $\mathbf{e}_\nu$. This choice avoids forcing the vehicle to chase replanned time-stamps when the tactical layer shifts the corridor.

The stage cost and the complete optimization problem are

$$\begin{aligned} \min_{\mathcal{U}_t} \quad & q_{\Delta\delta} (\delta_{0|t} - \delta_{-1|t})^2 + \sum_{k=0}^{N_c-1} [q_c e_{c,k+1|t}^2 + q_\psi e_{\psi,k+1|t}^2 + \mathbf{e}_{\nu,k|t}^\top R_\nu \mathbf{e}_{\nu,k|t}] \\ \text{s.t.} \quad & \boldsymbol{\chi}_{0|t} = \mathbf{0}, \\ & \boldsymbol{\chi}_{k+1|t} = f_d^{\text{kin}}(\boldsymbol{\chi}_{k|t}, \boldsymbol{\nu}_{k|t}), \quad k = 0, \dots, N_c-1, \\ & V_{\min} \leq V_{k|t} \leq v_{x,\max}, \\ & |\delta_{k|t}| \leq \delta_{\max}, \\ & |V_{k|t} - V_{k-1|t}| \leq \Delta V_{\max}, \\ & |\delta_{k|t} - \delta_{k-1|t}| \leq \dot{\delta}_{\max} T_c, \end{aligned} \quad (103)$$

where $R_\nu = \text{diag}(r_V, r_\delta)$, and $f_d^{\text{kin}}(\cdot)$ denotes the kinematic discrete dynamics in (95)–(96). The weight $q_{\Delta\delta}$ penalizes the first steering increment to enforce actuator continuity across cycles. The speed and steering rate constraints prevent saturation and ensure smooth transitions during corridor-side switches initiated by the tactical layer.

The NLP is solved in condensed form with $2N_c$ decision variables by a warm-started interior-point method. The receding-horizon policy applies only $(V_{0|t}^*, \delta_{0|t}^*)$ and re-solves at the next cycle. The actuator mapping converts this pair into the physical command vector

$$\mathbf{u}_{\text{act}} = \mathcal{M}(V_{0|t}^*, \delta_{0|t}^*) = [T_{\text{cmd}}, P_{b,\text{cmd}}, S_{\text{cmd}}, g_{\text{cmd}}]^\top, \quad (104)$$

where $T_{\text{cmd}} \in [0, 100]$ is the throttle percentage, $P_{b,\text{cmd}} \in [0, P_{b,\text{max}}]$ is the brake pressure, $S_{\text{cmd}} \in [S_{\text{min}}, S_{\text{max}}]$ is the steering actuator position, and $g_{\text{cmd}} \in \{1, \dots, G\}$ is the gear. The longitudinal force request driving the throttle/brake split is

$$F_x^{\text{req}} = \underbrace{c_{d,2}V^2 + c_{d,1}V + c_{d,0} + m \frac{V_{\text{la}} - V_{0|t}^*}{T_{\text{la}}}}_{F_{\text{ff}}} + \underbrace{K_p^V (V_{0|t}^* - V_{\text{meas}}) + K_i^V I_V}_{F_{\text{fb}}}, \quad (105)$$

and the allocation rule is

$$T_{\text{cmd}} = \begin{cases} \min\{100 F_x^{\text{req}}/F_{\text{max}}(\omega_e, g), T_{\text{tc}}\}, & F_x^{\text{req}} > 0, \\ 0, & \text{otherwise,} \end{cases} \quad (106)$$

$$P_{b,\text{cmd}} = \begin{cases} 0, & F_x^{\text{req}} > 0, \\ \text{clip}\left(\frac{-F_x^{\text{req}} r_w}{2k_b}, 0, P_{b,\text{max}}\right), & \text{otherwise,} \end{cases} \quad (107)$$

where T_{tc} is the traction-control limit and k_b converts pressure to wheel torque. The steering actuator is mapped as $S_{\text{cmd}} = r_s(\delta_{0|t}^* + \delta_{\text{off}})$ with calibrated ratio r_s and offset δ_{off} .

6 Simulation Results

6.1 High-Fidelity A2RL Simulation Environment

The closed-loop experiments were carried out in the high-fidelity simulator provided by the Abu Dhabi Autonomous Racing League (A2RL) [37], an international competition for full-scale autonomous racecars, which embeds a detailed Dallara EAV24 vehicle-dynamics model together with the Yas Marina track and environment model. Internally, the simulator and the decision stack communicate over a ROS2 platform, while the final actuator commands are delivered through a virtual CAN channel that mirrors the real race-vehicle interface. Sensor states, opponent states, and race-control flags are received from the simulator, and throttle, brake, steering, and gear commands are returned through the same loop. Because the decision stack runs as an external real-time process and all timing must survive the message transport layer, failures caused by delayed tactical updates, infeasible planner outputs, or controller fallback appear in the closed loop exactly as they would on the race vehicle.




Figure 3: A2RL high-fidelity simulation environment.

6.2 Parameter Configuration and Scenario Design

All numerical parameters used in the experiments are summarized in Table 1. The table intentionally reports only symbols and values so that the implementation settings can be inspected without interrupting the method derivations in the preceding sections.

Three overtaking scenarios are designed to cover complementary operating regimes that a competitive racing system must handle:

- **Scenario I (High-Speed Straight):** A pure straight section ($\kappa \approx 0$, minimum radius > 900 m) where both vehicles operate above 68 m/s. Tests whether the hierarchy can exploit a short high-speed closure window where decision latency directly determines success or failure.
- **Scenario II (Sharp-Corner):** A tight cornering section with minimum radius 10 m ($\kappa_{\max} = 0.098 \text{ m}^{-1}$) and speed variation from 12.5 to 65.8 m/s. Tests whether the corridor remains feasible through severe curvature where the GGV envelope restricts available acceleration and lateral space is tightest.
- **Scenario III (Mixed Straight-Corner):** A mixed section combining straights and corners ($\kappa_{\max} = 0.037 \text{ m}^{-1}$) against an equal-speed opponent. Tests the full adversarial capability where the ego must generate its own advantage from track-geometry exploitation.

Table 1: Parameter setting summary.

Sym.	Val.	Sym.	Val.	Sym.	Val.	Sym.	Val.	Sym.	Val.
m	742 kg	l_f, l_r	1.75, 1.35 m	I_z	1500 kg m ²	C_f, C_r	2.30, 2.40×10^5 N/rad	μ	1.20
$\delta_{\max}$	0.342 rad	$\dot{\delta}_{\max}$	0.5 rad/s	$[\underline{a}, \bar{a}]$	$[-5, 8] \text{ m/s}^2$	$v_{x,\max}$	90 m/s	ε_v	0.5 m/s
$c_{d,2}$	$-1.04765 \text{ N s}^2/\text{m}^2$	$c_{d,1}$	47.5 N s/m	$c_{d,0}$	-1609.73 N	p	0.75	T_s	0.05 s
V_{char}	80 m/s	n_{char}	8 m	a_{char}	20 m/s ²	Δs_{char}	50 m	$\Delta V_{s,\text{char}}$	30 m/s
g_{char}	50 m	$\tau_{\text{tc},\max}$	30 s	ε_τ	0.01 m/s	L_t^{cap}	[1.5, 6.0]	R_t^{cap}	[1.5, 6.0]
d_{safe}	[0.2, 0.8]	V_{cap}	[15, 90]	g_{chase}	[18, 24]	g_{ot}	[16, 24]	g_{abort}	[28, 45]
b_{lat}	$[-1.5, 1.5]$	w_{safe}	[1, 5]	q_{side}	$[-1, 1]$	$g_{\text{def}}^{\text{front}}$	25 m	$g_{\text{def}}^{\text{rear}}$	20 m
Δn_{safe}	2.5 m	$g_{\text{interlock}}$	6 m	T_{imm}	30 s	w_{ot}	30	w_{ovb}	-20
w_{col}	-100	w_{lat}	-2	w_{crush}	-5	w_{app}	3	w_{sm}	-0.1
w_{gs}	-0.5	w_{sbs}	0.5	w_{wide}	0.25	α_Q	3×10^{-4}	$ \mathcal{D} $	10^5
B	256	γ	0.99	τ	5×10^{-3}	f_{feat}	51 $\rightarrow$ 256 $\rightarrow$ 256	$f_{\pi,Q}$	[256, 256]
N_q	25	K	20	d_p	300 m	N_{sta}	150	Δs_{sta}	2 m
p_{fade}	3	S_{fade}	50 m	W_{fun}	2.5 m	W_{near}	5 m	λ_{side}	0.6
$L_{\text{min}}^{\text{corr}}$	3 m	β	0.2	k_{blend}	9	d_{rmp}	40 m	k_{sr}	15
w_ω	10	$w_{n,T}$	10	$w_{\xi,T}$	100	V_{min}	0.1 m/s	ε_κ	10^{-8}
$\varepsilon_n^{\text{sl}}$	0.1 m	$\varepsilon_\xi^{\text{sl}}$	0.1°	$N_{\text{iter}}^{\text{path}}$	5	q_n	50	q_ψ	5
q_v	1	R_a	0.05	R_δ	30	q_{prog}	0.5	w_ϵ	1000
N_c	15	T_c	0.05 s	$T_c N_c$	0.75 s				

6.3 Scenario I: High-Speed Straight

The ego vehicle enters the engagement with a large closing-speed advantage and must convert this advantage into a feasible lateral corridor before reaching the opponent. The run completes the pass with $\min \Delta s = -0.48$ m, $\Delta t_{\text{pass}} = 1.13$ s after entering OVERTAKE, $\max V = 73.10$ m/s, $\text{RMS}(e_n) = 0.02$ m, and $\text{RMS}(e_v) = 0.26$ m/s.




Figure 4: Scenario I engagement snapshots. The four representative instants show the planned tactical corridor and the local trajectory during approach, shadowing, commitment, and pass completion.

Figure 4 shows that the tactical decision is made before the vehicles become geometrically interlocked. In the first snapshot, the ego vehicle remains on a high-speed closure trajectory while the planner keeps the corridor centered enough to preserve recoverability. The second snapshot corresponds to the SHADOW phase: the tactical layer guides the ego vehicle toward the opponent's rear-quarter on the selected pass side, biasing the corridor asymmetrically to pre-position for the upcoming overtake while retaining a conservative fallback around the raceline. In the third snapshot, the FSM has committed to OVERTAKE; the corridor carver excludes the opponent footprint and forces the SOCP path to occupy the free side of the track. The final snapshot shows the pass completion, where the planned path has cleared the opponent without a discontinuous lateral jump. This progression is important because the learned component never commands steering or acceleration directly; it only reshapes the feasible set that the deterministic planner and MPCC must execute.




Figure 5: Scenario I closed-loop execution states.

Figure 5 quantifies the dynamic execution of the short pass. The path plot confirms that the ego trajectory remains inside the tactical corridor during the entire engagement, so the learned side decision does not create a geometric constraint violation. The speed trace reaches 73.10 m/s, indicating that the planner preserves the fast-closure advantage rather than sacrificing it through

excessive caution. The longitudinal and lateral acceleration histories remain smooth during the transition into OVERTAKE, which shows that the GGV-constrained speed profiler and MPCC friction slack prevent the maneuver from becoming a bang–bang lateral escape. The heading-error response stays bounded while the lateral path changes side, demonstrating that the corridor transition is converted into a controllable yaw demand rather than an abrupt heading correction.




Figure 6: Scenario I tracking and tactical-state responses.

Figure 6 shows that the controller tracks the planner output with sub-decimeter accuracy despite the very short decision window. The lateral error has an RMS value of only 0.02 m, so the vehicle remains effectively on the planned path during the pass. The speed tracking error is also small, with $\text{RMS}(e_v) = 0.26$ m/s, which confirms that the velocity cap and the GGV sweep are consistent with the MPCC’s achievable acceleration envelope. The FSM mode plot shows a clean transition from SHADOW—where the ego positions itself at the opponent’s rear-quarter—to OVERTAKE, while the corridor plot shows that the committed side is held during the pass rather than oscillating between alternatives. This stable side commitment is the key reason the planner can exploit the brief overtake opportunity without inducing high-frequency replanning.




Figure 7: Scenario I safety and performance envelope.

Figure 7(a) plots the longitudinal gap Δs and lateral gap Δn over time. Collision requires both $|\Delta s| < L_v$ and $|\Delta n| < W_v$ simultaneously; the two curves never enter their respective danger zones at the same instant, confirming collision-free execution. Figure 7(b) shows the 3-D GGV diagram where the scatter points represent actual accelerations during the engagement and the translucent surface is the speed-dependent handling envelope. On this high-speed straight, the vehicle operates

well within the GGV boundary as the near-zero curvature does not demand significant lateral acceleration.

6.4 Scenario II: Sharp-Corner

The ego vehicle must remain side-by-side with the opponent through a high-curvature sequence where the GGV envelope severely limits available acceleration. The run completes the pass with $\min \Delta s = -0.84$ m, $\Delta t_{\text{pass}} = 8.63$ s, $\max V = 65.83$ m/s, $\text{RMS}(e_n) = 0.05$ m, and $\text{RMS}(e_v) = 1.82$ m/s.




Figure 8: Scenario II engagement snapshots. The four representative instants show the planned tactical corridor and the local trajectory during the extended side-by-side overtake.

Figure 8 illustrates a qualitatively different engagement from Scenario I. The first instant shows the ego vehicle closing from behind while the corridor remains broad enough to retain both tactical options. The second instant captures the SHADOW phase: the ego vehicle is guided toward the opponent's rear-quarter on the chosen pass side, with the corridor progressively narrowing to steer the planner into the optimal lateral position for the upcoming side-by-side maneuver. During the third instant, the FSM transitions to OVERTAKE and the vehicles are side-by-side; the corridor excludes the opponent while preserving a continuous tube for the SOCP path. The fourth instant shows the ego vehicle ahead of the opponent, with the tactical layer relaxing the corridor after the pass. The sequence demonstrates that the same hierarchy handles not only short impulse-like passes, but also prolonged interactions where the corridor must remain stable for several seconds.




Figure 9: Scenario II closed-loop execution states.

Figure 9 shows that the longer engagement is executed with a lower peak speed of 65.83 m/s, reflecting the planner's need to remain in a viable side-by-side corridor rather than complete an immediate speed-dominant pass. The path remains inside the carved corridor, which confirms that the tactical layer maintains a feasible spatial tube even when the opponent occupies the critical region for a longer time. The acceleration plot shows larger dynamic transients than Scenario I, as expected for an extended interaction, but the response remains smooth enough for the MPCC to track without friction-limit divergence. The heading response remains bounded during the side-by-side phase, indicating that the SOCP path regularization prevents the lateral maneuver from degenerating into a sharp steering correction.




Figure 10: Scenario II tracking and tactical-state responses.

Figure 10 confirms that the increased duration of the duel does not destabilize the hierarchy. The lateral RMS tracking error is 0.05 m, still well below the corridor width and therefore small relative to the available passing margin. The speed error increases to 1.82 m/s because the controller must reconcile the speed profile with a longer side-by-side geometry, but the error remains bounded and does not prevent pass completion. The FSM plot shows a persistent OVERTAKE tenure rather than rapid mode switching, and the corridor plot shows a stable committed side throughout the decisive interval. This is the desired behavior for a learned tactical layer: once the policy selects a pass side, the downstream solvers receive a consistent feasible set long enough to execute the maneuver.




Figure 11: Scenario II safety and performance envelope.

Figure 11(a) confirms collision-free execution during the prolonged side-by-side phase: although Δs remains near zero for 8.6s, the lateral gap Δn stays well outside the collision zone throughout. Figure 11(b) shows that the sharp-corner overtake pushes the vehicle close to the GGV envelope boundary, with lateral accelerations reaching the handling limit at low speeds, confirming that the planner fully exploits the available tire grip to maintain the side-by-side corridor.

6.5 Scenario III: Mixed Straight-Corner

The ego vehicle must generate its own speed advantage through corridor exploitation rather than relying on a pre-existing closure-speed gap. The analysis focuses on the decisive 29s engagement window during which the ego vehicle identifies, commits to, and executes the overtake. The run completes the pass with $\min \Delta s = -0.12$ m, $\Delta t_{\text{pass}} = 20.13$ s after entering SHADOW, $\max V = 72.92$ m/s, $\text{RMS}(e_n) = 0.19$ m, and $\text{RMS}(e_v) = 1.62$ m/s.




Figure 12: Scenario III engagement snapshots. The ego vehicle overtakes an equal-speed opponent through sustained corridor exploitation, requiring the FSM to progress through Shadow, Overtake, and Hold modes during the decisive engagement window.

Figure 12 reveals a fundamentally different tactical structure from the previous two scenarios.

Because the opponent travels at the same nominal speed as the ego vehicle, the tactical layer cannot rely on a simple fast-closure strategy. Instead, the learned policy must identify a section of track geometry where the raceline speed surplus creates a transient advantage, commit to the corridor early enough to exploit this window, and sustain the overtake commitment through the cornering sequence. The SHADOW phase lasts 5.0s, during which the ego vehicle is progressively steered into the opponent's rear-quarter on the committed side, building the lateral offset necessary to enter OVERTAKE with a feasible corridor. The Overtake execution then requires 8.8s to clear the opponent and transition to Hold.




Figure 13: Scenario III closed-loop execution states.

Figure 13 shows the execution dynamics over the 29s engagement. The path plot demonstrates that the ego trajectory commits to a lateral excursion within the corridor as the tactical layer exploits the cornering speed differential. The speed trace shows variation between 22–73 m/s, reflecting the cornering sections where the GGV profiler limits speed while the tactical layer maintains commitment to the overtake. The acceleration magnitudes reach 32.8 m/s² longitudinally and 21.4 m/s² laterally, confirming that the ego vehicle operates near the handling limits to generate the speed advantage needed against an equal-speed opponent.




Figure 14: Scenario III tracking and tactical-state responses.

Figure 14 shows the tracking and FSM behavior during the adversarial engagement. The lateral RMS error is 0.19 m, larger than the two previous cases but still well within the corridor margin, confirming that the MPCC successfully tracks aggressive path references during the sustained side-by-side phase. The speed tracking error is 1.62 m/s, reflecting the controller's effort to maintain the speed advantage generated by the GGV profiler against an equal-speed adversary. The FSM mode plot shows the progression from SHADOW—where the ego is guided to the

opponent's rear-quarter to establish lateral pre-positioning—through OVERTAKE to HOLD, with the mode commitment sustained throughout the decisive interval. The pass completion with only 0.12 m of longitudinal margin confirms that the ego vehicle generates just enough advantage through cornering speed exploitation—a characteristic signature of same-speed adversarial racing where large surpluses are unavailable.




Figure 15: Scenario III safety and performance envelope.

Figure 15(a) demonstrates collision-free overtaking against the equal-speed opponent: the longitudinal gap decreases over 29s and crosses zero only at the final pass moment, while the lateral clearance remains safely outside the collision zone throughout. Figure 15(b) reveals that this scenario demands the broadest GGV utilization—the vehicle exploits both high longitudinal acceleration on straights and significant lateral acceleration through corners, covering a wide region of the handling envelope to generate the speed advantage needed against an opponent of equal capability.

6.6 Simulation Summary

Beyond the three representative cases, the same frozen tactical policy, planner, and controller were evaluated over 1000 closed-loop laps in the A2RL simulator. No online learning, case-specific retuning, or controller weight adjustment was used during this validation. The overtake success rate of 94.7% (947 out of 1000 laps) indicates that the learned tactical layer generalizes across a large number of repeated race interactions while preserving compatibility with the model-based planner and controller. Equally important, the validation was performed through the virtual-CAN simulation interface, so successful laps reflect end-to-end timing consistency rather than offline trajectory feasibility alone.

Overall, the experiments support the central design principle of the paper: learning is most effective when it determines the tactical feasible set, while the planner and controller enforce geometric consistency, handling-limit speed selection, and closed-loop vehicle execution. The short-pass, long-pass, and equal-speed adversarial cases show that the architecture is not tuned to a single overtaking geometry, and the 1000-lap validation demonstrates that this separation remains robust under repeated high-fidelity simulation.

6.7 Comparative Evaluation

To validate the proposed hierarchical game-guided corridor planning (HG-Corridor, ours), we compare it against four baselines under the same semi-physical simulation environment. All methods share the identical SOCP path planner and MPCC controller; only the tactical decision layer differs, thereby isolating the contribution of the learned corridor synthesis.

- **Pure Raceline:** No tactical layer. The planner follows the time-optimal raceline with no opponent awareness. Overtakes occur only when the ego vehicle's speed surplus naturally creates clearance.

- **A2RL OCP:** The official A2RL obstacle carver, which treats opponents as dynamic constraints in a single-shot optimal control problem without game-theoretic reasoning or mode commitment.
- **Rule-based GT:** A hand-tuned finite-state machine with fixed-threshold transitions ($g_{\text{chase}} = 20$ m, $g_{\text{ot}} = 8$ m, $g_{\text{abort}} = 35$ m). Represents the deterministic-rules approach common in autonomous racing.
- **Reactive APF:** An artificial potential field method that generates repulsive forces from opponents and attractive forces toward the raceline. Purely reactive—no strategic planning or mode persistence.

Table 2 reports the results averaged over three complementary scenarios (high-speed straight, sharp corners, and a full-lap engagement) with 1000 closed-loop laps per method. Note that the Reactive APF never attempts an overtake ($\text{OSR} = 0$), so the time-to-overtake (TTO) and average time-to-collision (Avg TTC) metrics are undefined and marked as “—”. This reflects the fundamental limitation of purely repulsive strategies: they guarantee safety by construction but sacrifice all competitive capability.

Table 2: Quantitative comparison against baseline algorithms over 1000 evaluation laps in the semi-physical A2RL simulator. Best values per metric are shown in **bold**. $\uparrow$: higher is better; $\downarrow$: lower is better.

Metric	HG-Corridor (Ours)	Pure Raceline	A2RL OCP	Rule-based GT	Reactive APF
OSR [%] $\uparrow$	94.7	38.2	22.2	51.6	0.0
TTO [s] $\downarrow$	4.82	7.25	12.50	9.37	—
Min Gap [m] $\uparrow$	5.83	1.42	0.10	2.14	12.3
Collisions $\downarrow$	0	21	19	48	0
Avg TTC [s] $\uparrow$	8.52	3.21	2.83	5.64	—
Track Violation [%] $\downarrow$	0.0	0.0	0.0	0.0	0.0
PFR $\uparrow$	1.000	0.988	1.000	0.987	0.985




Figure 16: Per-lap distribution of key safety and performance metrics across methods. Box plots show the median, interquartile range, and whiskers ($1.5 \times \text{IQR}$). HG-Corridor achieves both the tightest TTO distribution and the largest minimum gap, confirming consistent performance across laps. Reactive APF does not attempt overtakes (TTO, Avg TTC undefined).

HG-Corridor achieves an overtake success rate of 94.7%, nearly doubling the best baseline (Rule-based GT at 51.6%) while maintaining zero collisions and a minimum passing gap of 5.83 m. The A2RL OCP baseline closes to within 0.1 m and accumulates 19 collision events because it treats the opponent as a static obstacle without anticipating the interaction dynamics. The Rule-based GT produces 48 collisions because its fixed thresholds cannot adapt to the diversity of overtaking geometries across different track sections—a limitation that the learned TQC policy overcomes by adjusting gap thresholds and corridor width continuously from the 51-dimensional observation.

6.8 Ablation Study

The proposed HG-Corridor architecture integrates three key components beyond the standard SOCP+MPCC pipeline: (i) a hysteretic finite-state machine (FSM) that provides mode

persistence, (ii) learned corridor parameters from the TQC policy, and (iii) the TQC algorithm itself (vs. standard SAC). The ablation study isolates the contribution of each component by removing it from the full system while keeping all other modules unchanged. All variants are evaluated over the same 1000 laps in the semi-physical simulator.

- **w/o FSM:** The hysteretic FSM is bypassed. The RL policy output directly drives the corridor carver at every step without mode locking. The policy can switch between Shadow, Overtake, and Hold freely each cycle. This tests whether mode persistence is necessary for stable corridor generation.
- **w/o Learned Corridor:** The ten adaptive corridor parameters output by the TQC policy (width caps $L^{\text{cap}}, R^{\text{cap}}$, safety scaling d_{safe} , lateral bias b_{lat} , etc.) are replaced with fixed nominal values ($W_{\text{corr}} = 3.0$ m, $d_{\text{safe}} = 0.5$ m, $b_{\text{lat}} = 0$). The FSM still operates but receives no learned adaptation. This tests whether the policy’s continuous output contributes beyond mode selection.
- **SAC instead of TQC:** The distributional TQC policy is replaced with a standard Soft Actor-Critic trained under identical hyperparameters, network architecture, replay buffer, and reward. This isolates whether the pessimistic quantile truncation improves safety in tail-risk overtaking scenarios.

Table 3: Ablation study results. Each variant removes one component from the full HG-Corridor method. All evaluated over 1000 laps.

Metric	HG-Corridor	w/o FSM	w/o Learned Corr.	SAC (vs. TQC)
OSR [%] $\uparrow$	94.7	82.3	76.1	88.5
TTO [s] $\downarrow$	4.82	6.71	8.94	5.53
Collisions $\downarrow$	0	12	31	5
Min Gap [m] $\uparrow$	5.83	2.17	1.05	3.94
Avg TTC [s] $\uparrow$	8.52	4.31	2.86	6.17
PFR $\uparrow$	1.000	0.974	0.961	0.992




Figure 17: Per-lap distribution comparison for the ablation study. Removing the learned corridor produces the widest spread and lowest median in both Min Gap and Avg TTC, indicating frequent near-miss events. The full HG-Corridor maintains a consistently high safety margin with minimal variance.

The ablation reveals a clear hierarchy of component importance:

Removing the learned corridor has the most severe impact. OSR drops from 94.7% to 76.1% (−19.6 pp) and collisions increase from 0 to 31. The fixed-width corridor cannot adapt to narrow track sections or anticipate opponent trajectory evolution, leading to frequent infeasible corridor states (PFR drops to 0.961) and near-misses (min gap 1.05 m, barely one vehicle width). This confirms that the policy’s continuous corridor-shaping output—not just the discrete mode decision—is essential for safe high-performance overtaking.

Removing the FSM reduces OSR to 82.3% (−12.4 pp) and produces 12 collisions. The mode switching rate increases 4.6× (from 0.019 to 0.087 Hz), indicating that without hysteretic locking the policy oscillates between Shadow and Overtake at mode boundaries. Each oscillation causes the

corridor to flip side, which the SOCP planner cannot track in a single planning cycle, producing transient infeasibilities and collisions at the exact moment of corridor reversal.

Replacing TQC with SAC achieves a respectable 88.5% OSR but introduces 5 collisions and reduces the minimum gap from 5.83 to 3.94 m. The SAC policy selects corridor parameters that are optimal in expectation but occasionally expose the ego vehicle to tail-risk geometries where the opponent's actual trajectory deviates from the predicted mean. TQC's pessimistic quantile truncation avoids these corridors entirely, producing zero collisions at a modest cost of 6.2 pp in OSR relative to SAC.

7 Conclusion

In this paper, a hierarchical game-guided tactical corridor planning and control architecture is proposed for wheel-to-wheel autonomous racing in high-speed, highly dynamic adversarial scenarios. The framework restricts reinforcement learning to the tactical feasible-set synthesis layer, where a map-conditioned Truncated Quantile Critics policy embeds fixed track priors to generate behavior decisions, overtaking-side selections, and the parameters of a time-varying Frenet corridor, while a hysteretic finite-state machine and a deterministic carver convert these into auditable spatial constraints. Within the corridor, a sparse second-order cone program produces curvature-regularized paths, a forward-backward sweep enforces the GGV handling envelope, and a model predictive contouring controller executes the reference at the vehicle handling limit. By separating tactical intent from physical execution, the learned policy determines where the vehicle is permitted to drive while the model-based layers guarantee that tire limits, actuator bounds, and closed-loop dynamics are never violated, yielding an adaptive yet interpretable solution for competitive autonomous racing. Future work will extend the framework to multi-opponent races with multi-agent game equilibria, incorporate probabilistic opponent prediction into the corridor constraints, and transfer the approach to controlled track tests on the Dallara EAV24 platform.

References

- [1] B. Paden, M. Čáp, S. Z. Yong, D. Yershov, and E. Frazzoli, "A survey of motion planning and control techniques for self-driving urban vehicles," *IEEE Transactions on Intelligent Vehicles*, vol. 1, no. 1, pp. 33–55, 2016.
- [2] D. González, J. Pérez, V. Milanés, and F. Nashashibi, "A review of motion planning techniques for automated vehicles," *IEEE Transactions on Intelligent Transportation Systems*, vol. 17, no. 4, pp. 1135–1145, 2016.
- [3] M. Werling, J. Ziegler, S. Kammel, and S. Thrun, "Optimal trajectory generation for dynamic street scenarios in a Frenet frame," in *Proc. IEEE International Conference on Robotics and Automation (ICRA)*, 2010, pp. 987–993.
- [4] J. Ziegler, P. Bender, M. Schreiber, H. Lategahn, T. Strauss, C. Stiller, T. Dang, U. Franke, N. Appenrodt, C. G. Keller, E. Kaus, R. G. Herrtwich, C. Rabe, D. Pfeiffer, F. Lindner, F. Stein, F. Erbs, M. Enzweiler, C. Knöppel, J. Hipp, M. Haueis, M. Trepte, C. Brenk, A. Tamke, M. Ghanaat, M. Braun, A. Joos, H. Fritz, H. Mock, M. Hein, and E. Zeeb, "Making Bertha drive—An autonomous journey on a historic route," *IEEE Intelligent Transportation Systems Magazine*, vol. 6, no. 2, pp. 8–20, 2014.
- [5] J. Betz, A. Wischniewski, A. Heilmeier, F. Nobis, L. Hermansdorfer, T. Stahl, T. Herrmann, and M. Lienkamp, "A software architecture for the dynamic path planning of an autonomous racecar at the limits of handling," in *Proc. IEEE Intelligent Vehicles Symposium (IV)*, 2019, pp. 1–8.
- [6] J. Kabzan, M. de la Iglesia Valls, V. J. F. Reijgwart, H. F. C. Hendrikx, C. Ehmke, M. Prajapat, A. Bühler, N. Gosala, M. Gupta, R. Sivanesan, A. Dhall, E. Chisari, N. Kumar, I. A. Paireda, and A. Liniger, "AMZ driverless: The full autonomous racing system," *Journal of Field Robotics*, vol. 37, no. 7, pp. 1267–1294, 2020.
- [7] J. Betz, H. Zheng, A. Liniger, U. Rosolia, P. Karle, M. Behl, V. Krovi, and R. Mangharam, "Autonomous vehicles on the edge: A survey on autonomous vehicle racing," *IEEE Open Journal of Intelligent Transportation Systems*, vol. 3, pp. 458–488, 2022.
- [8] J. Betz, T. Betz, F. Fent *et al.*, "TUM autonomous motorsport: An autonomous racing software for the Indy Autonomous Challenge," *Journal of Field Robotics*, vol. 40, no. 4, pp. 783–809, 2023.

- [9] C. Jung, A. Finazzi, H. Seong *et al.*, “An autonomous system for head-to-head race: Design, implementation and analysis,” *IEEE Transactions on Intelligent Transportation Systems*, vol. 24, no. 10, pp. 10470–10486, 2023.
- [10] A. Liniger and J. Lygeros, “A noncooperative game approach to autonomous racing,” *IEEE Transactions on Control Systems Technology*, vol. 28, no. 3, pp. 884–897, 2020.
- [11] M. Wang, Z. Wang, J. Talbot, J. C. Gerdes, and M. Schwager, “Game-theoretic planning for self-driving cars in multivehicle competitive scenarios,” *IEEE Transactions on Robotics*, vol. 37, no. 4, pp. 1313–1325, 2021.
- [12] D. Fridovich-Keil, E. Ratner, L. Peters, A. D. Dragan, and C. J. Tomlin, “Efficient iterative linear-quadratic approximations for nonlinear multi-player general-sum differential games,” in *Proc. IEEE International Conference on Robotics and Automation (ICRA)*, 2020, pp. 1475–1481.
- [13] F. Laine, D. Fridovich-Keil, C.-Y. Chiu, and C. J. Tomlin, “Multi-hypothesis interactions in game-theoretic motion planning,” in *Proc. IEEE International Conference on Robotics and Automation (ICRA)*, 2021, pp. 8016–8023.
- [14] R. Spica, E. Cristofalo, Z. Wang, E. Montijano, and M. Schwager, “A real-time game theoretic planner for autonomous two-player drone racing,” *IEEE Transactions on Robotics*, vol. 36, no. 5, pp. 1389–1403, 2020.
- [15] Y. Jia, M. Bhatt, and N. Mehr, “RAPID: Autonomous multi-agent racing using constrained potential dynamic games,” in *Proc. European Control Conference (ECC)*, 2023, pp. 1–8.
- [16] D. Kalaria, C. Maheshwari, and S. Sastry, “ α -RACER: Real-time algorithm for game-theoretic motion planning and control in autonomous racing using α -potential function,” in *Proc. Learning for Dynamics and Control Conference (L4DC)*, PMLR, vol. 283, pp. 1–18, 2025.
- [17] P. R. Wurman, S. Barrett, K. Kawamoto *et al.*, “Outracing champion Gran Turismo drivers with deep reinforcement learning,” *Nature*, vol. 602, no. 7896, pp. 223–228, 2022.
- [18] F. Fuchs, Y. Song, E. Kaufmann, D. Scaramuzza, and P. Dürri, “Super-human performance in Gran Turismo Sport using deep reinforcement learning,” *IEEE Robotics and Automation Letters*, vol. 6, no. 3, pp. 4257–4264, 2021.
- [19] Y. Song, H. Lin, E. Kaufmann, P. Dürri, and D. Scaramuzza, “Autonomous overtaking in Gran Turismo Sport using curriculum reinforcement learning,” in *Proc. IEEE International Conference on Robotics and Automation (ICRA)*, 2021, pp. 9403–9409.
- [20] E. Kaufmann, L. Bauersfeld, A. Loquercio, M. Müller, V. Koltun, and D. Scaramuzza, “Champion-level drone racing using deep reinforcement learning,” *Nature*, vol. 620, no. 7976, pp. 982–987, 2023.
- [21] H. Lin, Y. Song, and D. Scaramuzza, “Reaching the limit in autonomous racing: Optimal control versus reinforcement learning,” *Science Robotics*, vol. 8, no. 82, art. eadg1462, 2023.
- [22] F. Christ, A. Wischnewski, A. Heilmeier, and B. Lohmann, “Time-optimal trajectory planning for a race car considering variable tyre-road friction coefficients,” *Vehicle System Dynamics*, vol. 59, no. 4, pp. 588–612, 2021.
- [23] A. Heilmeier, A. Wischnewski, L. Hermansdorfer, J. Betz, M. Lienkamp, and B. Lohmann, “Minimum curvature trajectory planning and control for an autonomous race car,” *Vehicle System Dynamics*, vol. 58, no. 10, pp. 1497–1527, 2020.
- [24] N. R. Kapania, J. Subosits, and J. C. Gerdes, “A sequential two-step algorithm for fast generation of vehicle racing trajectories,” *ASME Journal of Dynamic Systems, Measurement, and Control*, vol. 138, no. 9, art. 091005, 2016.
- [25] T. Lipp and S. Boyd, “Minimum-time speed optimisation over a fixed path,” *International Journal of Control*, vol. 87, no. 6, pp. 1297–1311, 2014.
- [26] R. Verschuere, S. De Bruyne, M. Zanon, J. V. Frasch, and M. Diehl, “Towards time-optimal race car driving using nonlinear MPC in real-time,” in *Proc. IEEE Conference on Decision and Control (CDC)*, 2014, pp. 2505–2510.

- [27] S. He, J. Zeng, and K. Sreenath, “Autonomous racing with multiple vehicles using a parallelized optimization with safety guarantee using control barrier functions,” arXiv preprint arXiv:2112.06435, 2021.
- [28] A. Wischnewski, T. Herrmann, F. Werner, and B. Lohmann, “A Tube-MPC approach to autonomous multi-vehicle racing on high-speed ovals,” *IEEE Transactions on Intelligent Vehicles*, vol. 8, no. 1, pp. 368–378, 2023.
- [29] A. Liniger, A. Domahidi, and M. Morari, “Optimization-based autonomous racing of 1:43 scale RC cars,” *Optimal Control Applications and Methods*, vol. 36, no. 5, pp. 628–647, 2015.
- [30] U. Rosolia and F. Borrelli, “Learning how to autonomously race a car: A predictive control approach,” *IEEE Transactions on Control Systems Technology*, vol. 28, no. 6, pp. 2713–2719, 2020.
- [31] J. Kabzan, L. Hewing, A. Liniger, and M. N. Zeilinger, “Learning-based model predictive control for autonomous racing,” *IEEE Robotics and Automation Letters*, vol. 4, no. 4, pp. 3363–3370, 2019.
- [32] M. Brunner, U. Rosolia, J. Gonzales, and F. Borrelli, “Repetitive learning model predictive control: An autonomous racing example,” in *Proc. IEEE Conference on Decision and Control (CDC)*, 2017, pp. 2545–2550.
- [33] T. Novi, A. Liniger, R. Capitani, and C. Annicchiarico, “Real-time control for at-limit handling driving on a predefined path,” *Vehicle System Dynamics*, vol. 58, no. 7, pp. 1007–1036, 2020.
- [34] U. Rosolia and F. Borrelli, “Learning model predictive control for iterative tasks: A data-driven control framework,” *IEEE Transactions on Automatic Control*, vol. 63, no. 7, pp. 1883–1896, 2018.
- [35] L. Hewing, K. P. Wabersich, M. Menner, and M. N. Zeilinger, “Learning-based model predictive control: Toward safe learning in control,” *Annual Review of Control, Robotics, and Autonomous Systems*, vol. 3, pp. 269–296, 2020.
- [36] L. Brunke, M. Greeff, A. W. Hall *et al.*, “Safe learning in robotics: From learning-based control to safe reinforcement learning,” *Annual Review of Control, Robotics, and Autonomous Systems*, vol. 5, pp. 411–444, 2022.
- [37] Abu Dhabi Autonomous Racing League (A2RL), “A2RL: Autonomous racing competition,” Abu Dhabi, UAE, 2024. [Online]. Available: <https://a2rl.io>